



video search & epidemic protocols in Zero-Server P2P Systems



Dr. Ir. Johan Pouwelse
Assistant Professor
j.a.pouwelse@TUDelft.nl
Delft University of Technology






Contents

- Two different search approaches
- Query-response based systems
- Internet-deployments
 - Napster, Gnutella, Kazaa
- Academic research
- Epidemic protocols
 - Fundamentals
 - Examples
 - Buddycast case study
- P2P Video search: open research challenges


Two different search approaches



Proactive Pull



Background push (Epidemic protocols)





P2P keyword search approaches

- Central server, flooding, superpeers
- Semantic routing
- DHT
 - “keyword lookup can be layered on top of the distributed hash model, it is an open question if indexing can be done efficiently”*
- Cubit, 2008, edit distance.

**Source "Looking up data in P2P systems",
H Balakrishnan, F Kaashoek, D. Karger, R. Morris, I. Stoica,
Communications of the ACM, February 2003*

Kazaa Lite - [Search]

FileViewPlayerToolsActionsHelp

Web

My Kazaa Lite

Theater

Search

Traffic

Kazaa Lite.com Forums

New search

Download

Search Field

Search

Search for video files

Search for:

two towers

Found 231 files. Press Search More to get more results.

All

Title

Artist

Search More

Stop Search

« Back

More search options »

Search filter

Search Tip

Type in what you are looking for. If you are looking for a specific media type you may check the corresponding button to get more search options.

Title	User	Size
lords of the ring(3of3)	kazaa Lite...	1
[tmd]The.Lord.Of.The.Rings.The.Two.Towers.(NEWFXP).ts.(1of3)-real deal	(zenit)@...	1
Lord of the Rings The Two Towers (1of3)	kaak@K...	1
Lord Of The Rings: The Two Towers Preview (No Fake!) DIVX	2 Users	1
(tmV)Lord.Of.The.Rings.--.The.Two.Towers.(CD3of3)	3 Users	1
(tmV)Lord.Of.The.Rings.--.The.Two.Towers.(CD2of3)	2 Users	1
The Lord Of The Rings The Two Towers (2002) cd 1 of 3	6 Users	7
The Lord Of The Rings The Two Towers cd 3 of 3	9 Users	2
Lord Of The Rings(full)	4 Users	7
lord of the rings-the two towers cd2of3.mpg	2 Users	6
Lord of The Rings- The Two Towers (Part 2/2)	3 Users	1
[TMD]The.Lord.Of.The.Rings.The.Two.Towers.(TCF).INTERNAL.CAM.(3of3)	15 Users	1
The.Lord.Of.The.Rings.The.Two.Towers.INTERNAL.CAM-TCF.CD1	2 Users	6
the lord of the rings II	4 Users	7
The Lord Of The Rings - The Two Towers [INTERNAL.CAM-TCF.CD3]	2 Users	5
The Lord of the Rings - The Two Towers (2002 FULL MOVIE) (DivX5) (640x352)	9 Users	1
Lord Of The Rings The Two Towers cd 2	6 Users	7
The Lord Of The Rings The Two Towers (2002) cd 2 of 3	2 Users	7
Lord of the Rings- Two Towers FULL (1/2) [DivX]	6 Users	6
the_two_towers-trailor_fs.l	3 Users	4
Lord of the Rings: The Two Towers	13 Users	7
Two towers trailer	2 Users	3
Lord Of The Rings - The Two Towers Trailer Large	kvist@K...	3
(Smr)lord of the rings - The two towers (Real. Ripped From Studio;)) 1 of 3	5 Users	1
Sagan om de två tornen-Trailer	matzz@...	3
Lord of the Rings 2 - The Two Towers	4 Users	6

Play

Pause

Stop

Previous

Next


Progress

Volume

Repeat

My Participation Level: Supreme Being (Kazaa Lite.com) 4,043,950 users online, sharing 762,697,431 files (5,720,44E) Connected as synctext@KaZaA, sharing 10 files


Semantic routing



A query is forwarded to nodes that are likely to have knowledge about a reply based on semantics.


Explored in: Socialised.net, Neurogrid, Alpine, Limewire

Multi-keyword search using DHT and Bloom filters



- s_A and s_B : servers storing the document ID lists for keywords k_A and k_B .
- A and B are the sets of document IDs matching the keywords k_A and k_B .
- 3. $F(A)$ is a Bloom filter representation of A .

Cubit: approximate keyword search



1. Edit nearness
2. No semantics
3. Released in 2007

Source: <http://www.cs.cornell.edu/~bwong/Cubit>

Cubit search



Azureus

File View Tools Window Help

Vuze beta

Dashboard On Vuze Library

My Torrents Cubit Search

Search Options

Search String: Number of Results:

Filename	Comment	Torrent Link
Fedora-9-i686-Live	Fedora-9-i686-Live	http://torrent.fedor...
Fedora-7-Live-i686	Fedora-7-Live-i686	http://torrent.fedor...
Fedora-7-source	Fedora-7-source	http://torrent.fedor...
Fedora-9-x86_64-DVD	Fedora-9-x86_64-DVD	http://torrent.fedor...
Fedora-9-x86_64-Live-KDE	Fedora-9-x86_64-Live-KDE	http://torrent.fedor...
Fedora-8-Live-KDE-i686	Fedora-8-Live-KDE-i686	http://torrent.fedor...
Fedora-8-dvd-i386	Fedora-8-dvd-i386	http://torrent.fedor...



Epidemic protocol fundamentals

- Each peer has information regarding the state of the entire distributed system
- Each peer periodically chooses a peer with which to share this its information
- New information is merged with existing information to obtain a new state
- Design space: target peer, info exchange, retention
- Challenges
 - Availability
 - Connectivity
 - Security



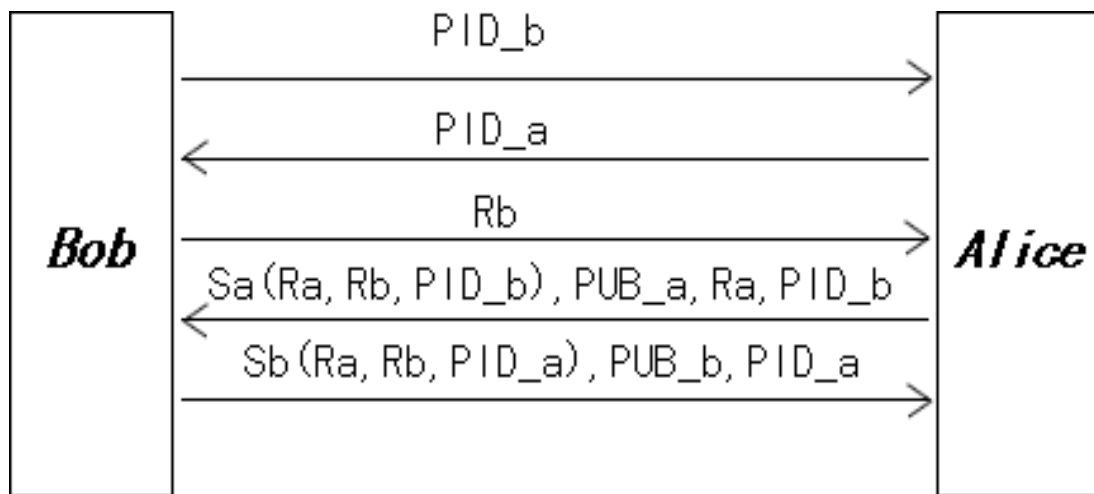
Epidemic protocol examples

- Xerox Clearinghouse service (1987); replica consistency
- Newscast (2002)*, vicinity, T-MAN, BAR gossip
- BuddyCast (2006); *first Internet-deployed protocol*
 - Limit randomness
 - Avoid one-time encounters
 - Both peer and content discovery
 - Apply collaborative filtering
 - Exploit the TeraByte hard-disk age
 - Preserve state across sessions


**Source: M. Jelasity, M. van Steen. "Large-Scale Newscast Computing on the Internet." Technical Report IR-503, October 2002.*

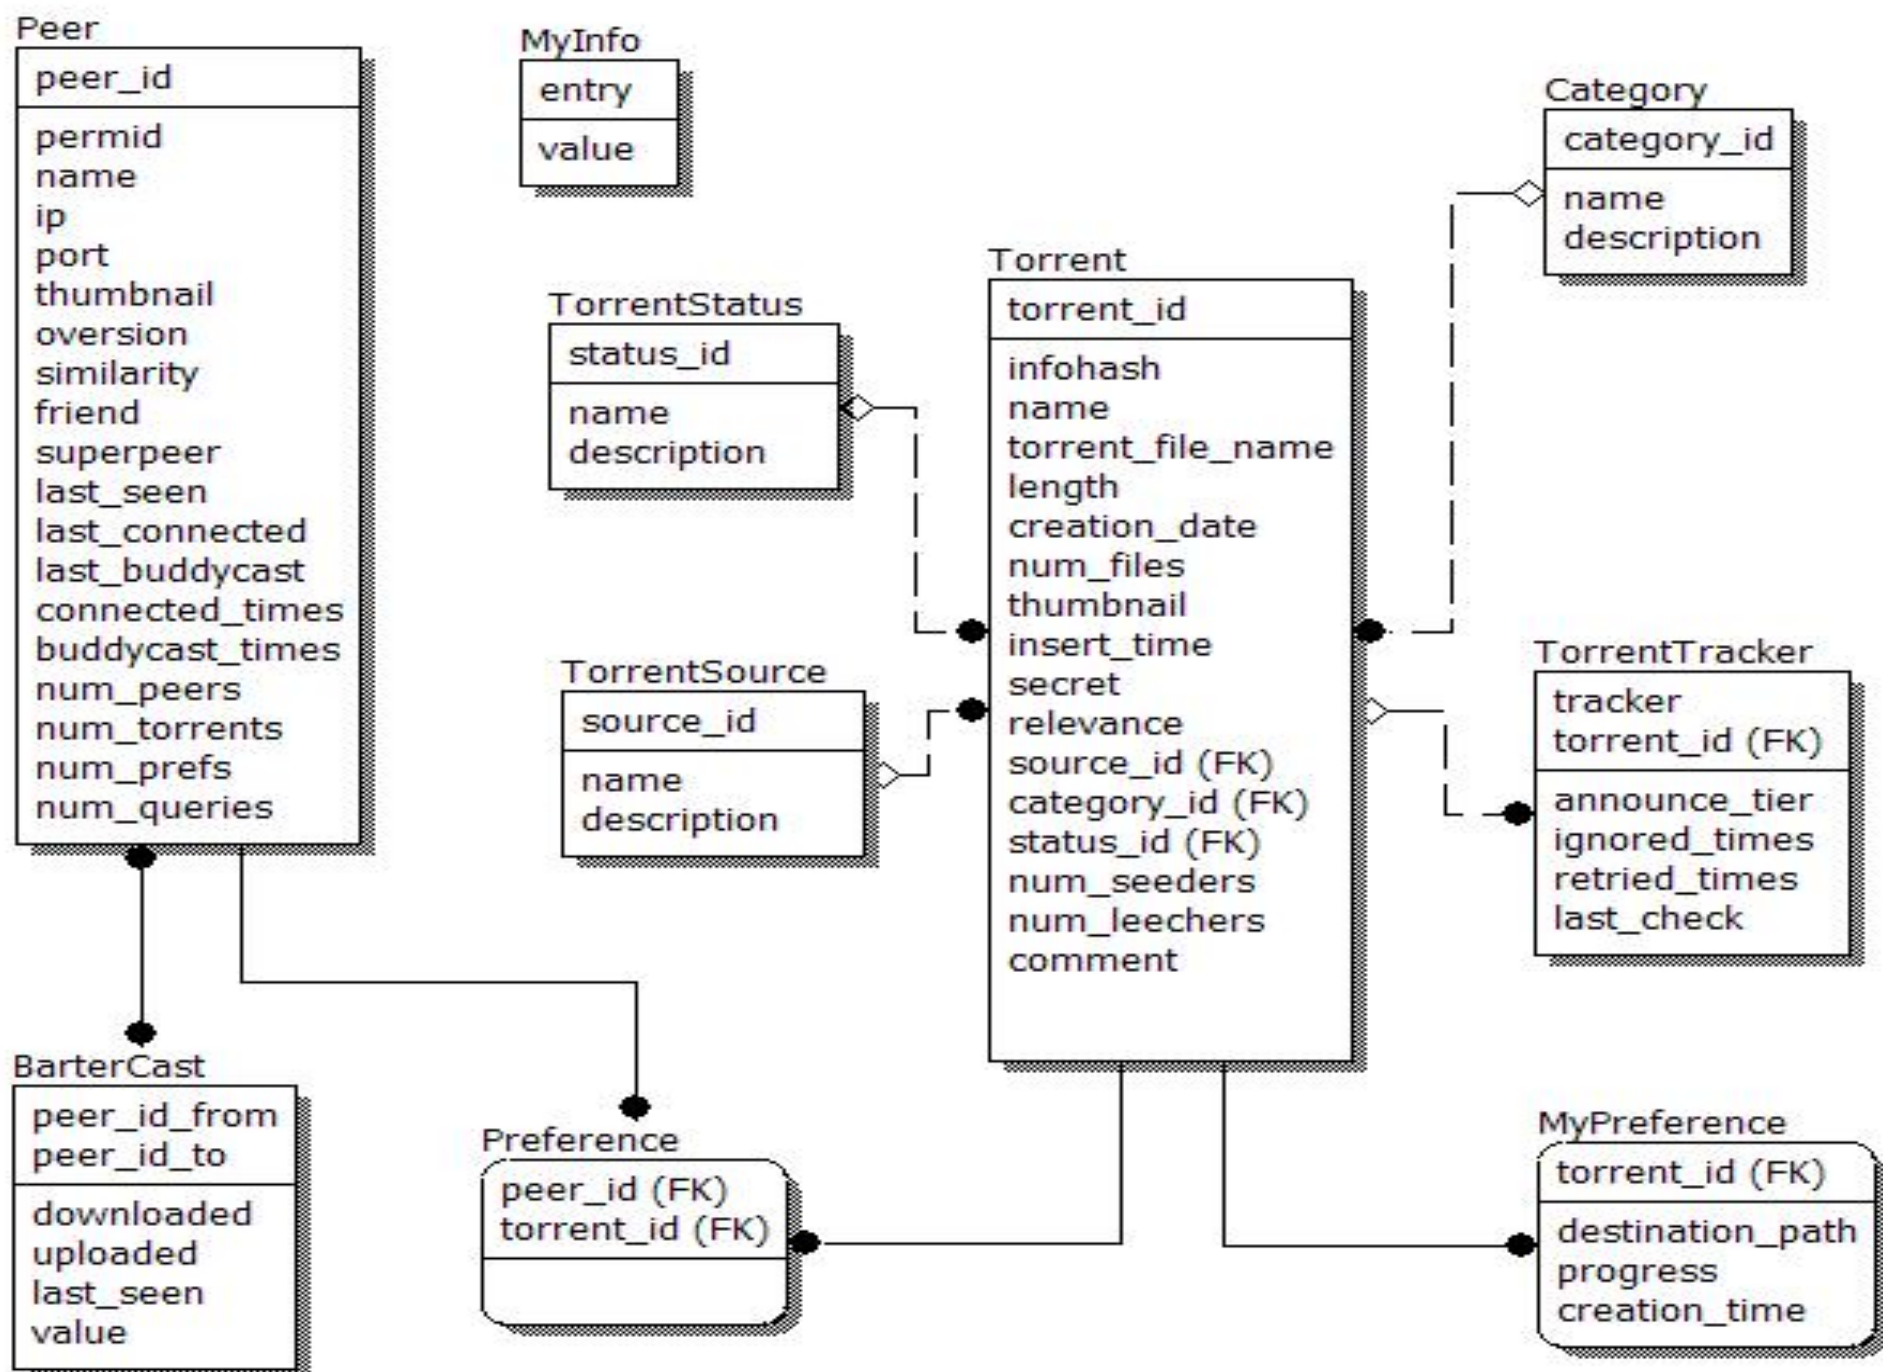
Preserve state across sessions

- Storage of state in relational database
- Long-term peer identifier (e.g. PermID)
- Robust against IPv4 changes
- PermID: Non-spoofable and compact
- Store:
content, peers, profiles, social network...




*R:Random number, S:Signature, PID:BT PeerID, PUB:PermID






Semantic clustering

- Organize P2P network
- References to related items
- Used in Buddycast
- Message layout
 - Your preferences
 - Your similar peers (*taste buddies*)
 - Their preferences
- Emergence
 - Peer discovery
 - Content discovery
 - Database of tastes



Semantic discovery



BuddyCast message: peers and Bittorrent items



Recent downloads (max 25)
{SHA1,seed/leech,last_seen}


Recent discovered (max 25)
{SHA1,seed/leech,last_seen}

Taste Buddy (max 10)
{IP,Port,Similarity,last_seen}

Random Peer (max 10)
{IP,Port,Similarity,last_seen}




BuddyCast trace of P2P downloads





Buddycast 1: no online check (2006)





Improvement: Keep 30 Open TCP connections

CC: Connection Candidate Cache (size=100)

Potential targets for next connection

CT: Connected Taste buddies

Most similar discovered peers (size=10)

CR: Connected Random peers

Last connected peers (size=10)

CU: Connected NATed/Firewalled peers

Connected peers with (size=10)

loop:

wait(ΔT time units) {wait=15 seconds}

$Q \leftarrow$ Top similar peer or likely online peer from CC

connectPeer(Q)

remove Q from CC

if Q is connected successfully then:

buddycast msg send \leftarrow createBuddycastMsg()

send buddycast msg send to Q


receive buddycast msg recv from Q

CC \leftarrow fillPeers(buddycast msg recv)

addConnectedPeer(Q) {add to CT, CR or CU using sim¹⁹()}



Buddycast 2: **with** online check



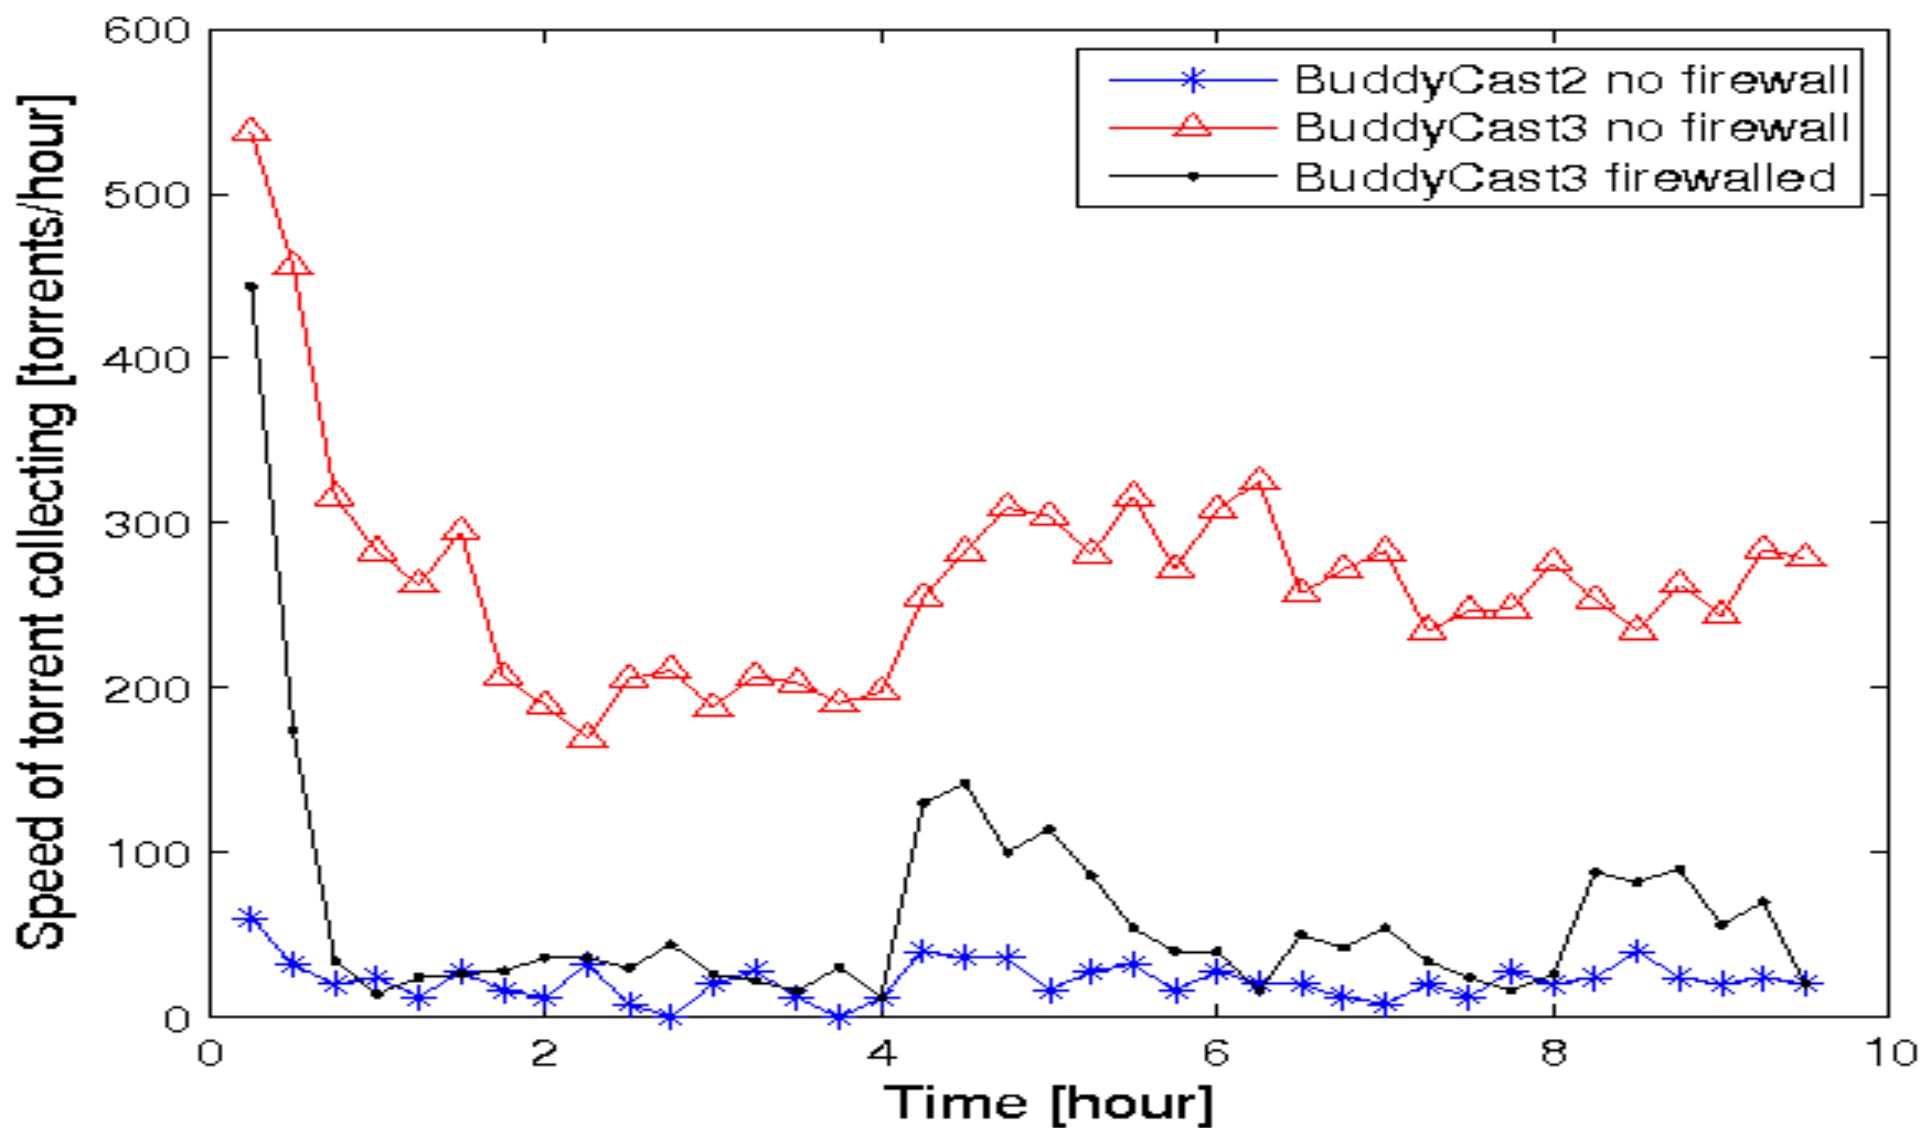


Epidemic protocol and keyword search

- Peers gather peers, items, profiles (size=100k)
- Keyword search
 - Local database
 - N peers around you (x 100K items)
"remote search"
 - Validate validity of P2P items
- Hybrid solution:
proactive pull + constant background push



Buddycast3: dynamic speed in 1st hour



files

persons

profile

library

subscr

friends

family filter: ON OFF



discovered files: 7258
 discovered persons: 3083
 Person connected:
 65.191.76.140:6881 36.765
 (peer deemed unreachable)

Discovered files

legal torrents

stop clear results

Searching for 'legal torrents'... (112 results)

compressed

video

[Exclusive] Tiesto - Club Life 026 09-28-07 -

TUTORIAL: How to use BitTorrent and download Torrents

Fairy From Wonderland Tian Wai Fei Xian episode 1

Top 100 Fonts (Graffiti, Horror, Handwritten More) -

DivX for Windows (incl. DivX Player) v6.6.2 -LegalTorrents

Fairy From Wonderland -Tian Wai Fei Xian episode

Free Legal Movies Videos To Download?

Cool Youtube Downloader Full - LegalTorrents

PeerIt.com Peer to Peer Signup and Sell Example

Fairy From Wonderland -Tian Wai Fei Xian episode

Fairy From Wonderland Tian Wai Fei Xian episode 1

DVD Shrink v3.2.0.15 - LegalTorrents

Mac Tips - Transmission

Fairy From Wonderland Tian Wai Fei Xian episode 1

Freeware Is Better Than Torrents - Duncan's Freeware

The Complete Works of William Shakespeare eBook

What is a Torrent? Is BitTorrent Legal?

Fairy From Wonderland Tian Wai Fei Xian episode 1

The Art of Writing Speaking the English Language eBook

[www.torsky.org] Jassie from Barely Legal

other

other

other

other

1 2 3 4

Fairy From Wonderland Tian Wai Fei Xian

options

info

files

size: 11:43

creation 23-10-2006

popularity: 15385

fit to taste: ♥

Description

from youtube: the first episode of Fairy From Wonderland or Tian Wai Fei Xian with english subs-----; dont think this is illegal t...

People who like this also like

No similar files found yet.

Similar Files

No similar files found yet.



Video search: open research challenges

- Spam & fraud prevention
- Rich metadata: tagging and ratings, subtitles, annotations..
- Moderation
- Relevance ranking
 - Semantic matching; word boundary, swarm or file name..
 - Popularity and age
 - Perceptual quality and download speed
 - Bonus or penalty based on collaborative filtering
 - Implicit feedback usage; Clicklogs



Final remarks

- What is P2P Video without search capability?
- Maturing field and transition:
 - From simulations to prototypes
 - From models to datasets
- Deployed self-organizing systems
- Importance of on-line video
 - Youtube influences wars and elections
 - Media is where society thinks
 - Mediaspace controlled by the people
- New business opportunities: reach millions at near-zero cost 25